

Evaluating the Nix Evaluator

Why Nix Performance Sometimes... Doesn't



Connor Baker

2025-03-07

Planet Nix

Topics covered

- Benchmarking setup
- Nix evaluation performance over time
- Suggested areas for improvement

Assumptions

1. Can improve?

1. Can improve?
 - Historically, yes!

1. Can improve?

- Historically, yes!

2. Should improve?

1. Can improve?

- Historically, yes!

2. Should improve?

- It depends!

Benchmarking

Benchmarking is **difficult**.

What can we easily measure?

- Data reported by `NIX_SHOW_STATS`
 - CPU/GC time, number of certain operations, etc.
- Data reported by `GNU time`
 - IO: context switches, page faults, etc.
 - Memory: page size, maximum resident set size, etc.
 - Time: real, user, and sys time

- A Nix flake for benchmarking the Nix flake¹
- Matrix Nix packages and configurations through flakes
- Runs `time nix eval` inside the sandbox n times
- Collects the results with some additional metadata into JSON
- Data is suitable for visualization with VegaLite

¹<https://github.com/ConnorBaker/benchmarking-nix-eval>

This presentation uses **VegaLite**
through **WASM** as a **Typst** package.

Examples

Testbed setup

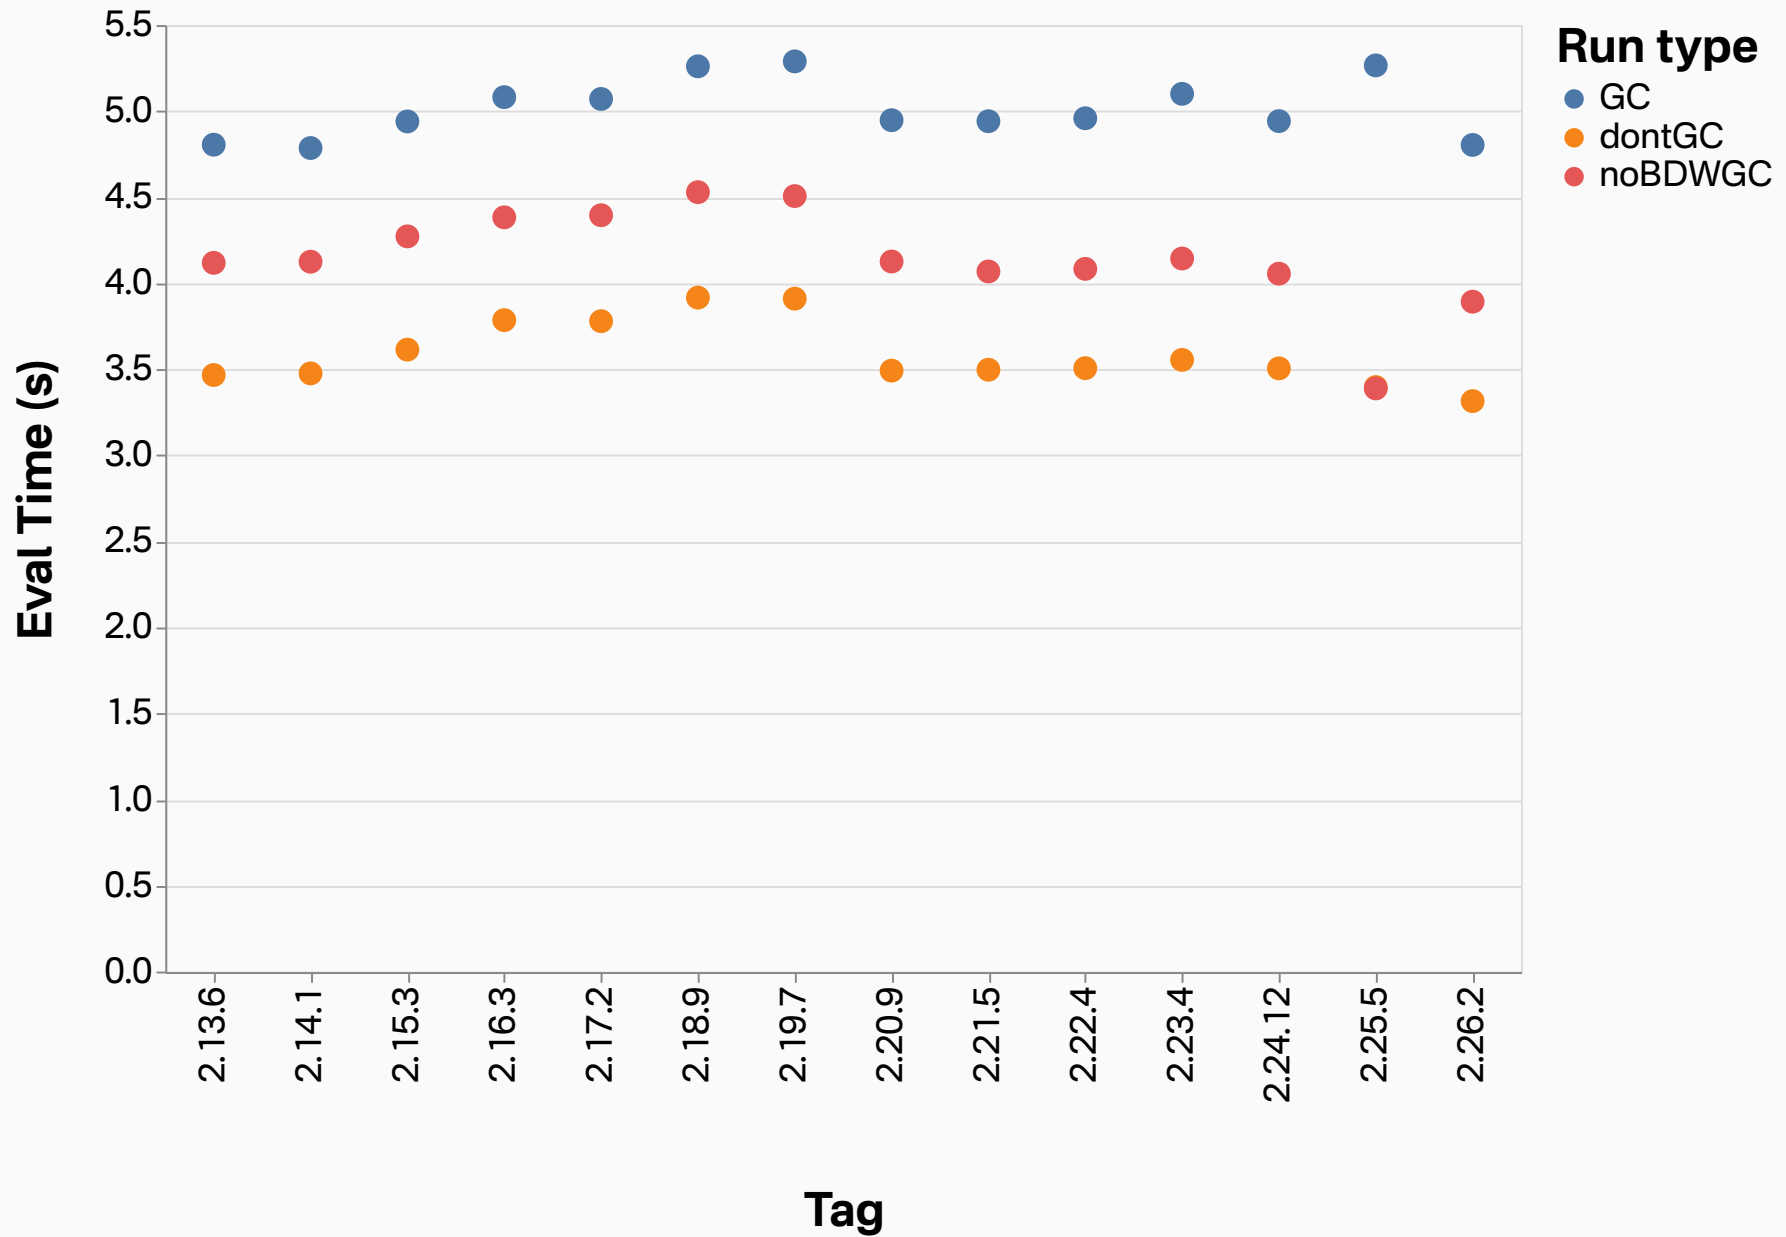
- Intel i9-13900K @ 3 GHz
 - Did not change niceness/pin to a favored core
- 96 GB DDR5 RAM
 - Did not attempt flushing caches
- Four-way ZFS RAID0
 - No deduplication/compression/integrity checking (just ARC)
 - Did not change IO niceness/flush caches
- Linux 6.12.13
- NixOS unstable @ 2ff53fe (2025-02-13)
- `mimalloc` as the default allocator

Software setup


- Latest minor versions of Nix (2.13-2.26)
- Benchmarks run one at a time, 20 times each for each config
 - With collection (GC)
 - Without collection (dontGC)
 - Without BDWGC (noBDWGC)
- Median values are plotted
 - Observed little variation between runs
- Generated data is available¹

¹<https://github.com/ConnorBaker/benchmarking-nix-eval/releases/download/v0.0.1/aggregated-nixos-desktop-20-runs-1-job-no-boost.json>


firefox-unwrapped eval time




firefox-unwrapped eval space




release-attrpaths-superset.names eval time




release-attrpaths-superset.names eval space




closures.smallContainer.x86_64-linux eval time




closures.smallContainer.x86_64-linux eval space




closures.lapp.x86_64-linux eval time




closures.lapp.x86_64-linux eval space




closures.kde.x86_64-linux eval time




closures.kde.x86_64-linux eval space



iso_gnome.x86_64-linux eval time



iso_gnome.x86_64-linux eval space



Summary

- If you need faster evaluation, set `GC_DONT_GC`
 - `nix-eval-jobs` (and Hydra) do this¹
- No GC is slower than using BDWGC
 - Individual allocations vs. batched allocations
- No GC uses less memory than BDWGC
 - No bookkeeping overhead
- Evaluation should be separate from builds
 - Cache derivations
 - Build separately to avoid resource contention

¹<https://github.com/nix-community/nix-eval-jobs/blob/4b392b284877d203ae262e16af269f702df036bc/src/nix-eval-jobs.cc#L421-L422>

Evaluator structures

- Less than 20 possible (internal) types
- Structure is 24 bytes
 - 8 bytes (due to padding) for the type
 - 16 bytes for the actual content
- Created everywhere during evaluation

- 0/1/2 element lists are inlined into a `Value`
- Otherwise, a C-style array of `Value` *
- Fantastic data locality
- No sharing of existing values

- C-style array of `Attr`, a structure with three fields
 - `Symbol name` (4 bytes)
 - `PosIdx pos` (4 bytes)
 - `Value * value` (8 bytes)

- Primitives for operating on values should:
 - be performant
 - be composable
- Data structure implementation should make the primitives exposed

Improvements

Suggested improvements should be **orthogonal** to those an **optimizing** or **parallel interpreter** would provide.

Data structures supporting **sharing**.

Shrinking the **Value** struct.

Future work

Future work

- Modularizing benchmarking-nix-eval
- Adding more benchmarks
- Building a web dashboard to visualize the data
- Integration into CI to detect regressions